

Model Developed by

Name: **Umair Ali**

Contact: **+923480233673**

Email: **uape00@gmail.com**

GitHub: <https://github.com/1umairali/models>

Developed a machine learning-based system for detecting brain tumors from MRI images. The project involved preprocessing image data and extracting relevant features, followed by training and evaluating various machine learning classifiers, including Support Vector Classifier (SVC), Logistic Regression, k-Nearest Neighbors (k-NN), Naives Regression, Random Forest, and Decision Tree. The models were assessed based on accuracy and performance metrics to determine the most effective approach for tumor classification. Demonstrated strong skills in medical image analysis and classical machine learning techniques using Python and scikit-learn.

```
In [1]: import numpy as np # for numeric calculation
import pandas as pd # for data analysis and manipulation
import matplotlib.pyplot as plt # for data visualization
import seaborn as sns # for data visualization
```

```
In [2]: # import dataframe
url = 'https://raw.githubusercontent.com/1umairali/models/main/brain_tumor_detection/brain_dataset.csv'
brain_dataframe = pd.read_csv(url)
brain_dataframe
```

Out[2]:

| | Image | Class | Mean | Variance | Standard Deviation | Entropy | Skewness | Kurtosis | Contrast | Energy | ASM | Homogeneity |
|-------------|-----------|-------|-----------|-------------|--------------------|----------|----------|-----------|------------|----------|----------|-------------|
| 0 | Image1 | 0 | 6.535339 | 619.587845 | 24.891522 | 0.109059 | 4.276477 | 18.900575 | 98.613971 | 0.293314 | 0.086033 | 0.530941 |
| 1 | Image2 | 0 | 8.749969 | 805.957634 | 28.389393 | 0.266538 | 3.718116 | 14.464618 | 63.858816 | 0.475051 | 0.225674 | 0.651352 |
| 2 | Image3 | 1 | 7.341095 | 1143.808219 | 33.820234 | 0.001467 | 5.061750 | 26.479563 | 81.867206 | 0.031917 | 0.001019 | 0.268275 |
| 3 | Image4 | 1 | 5.958145 | 959.711985 | 30.979219 | 0.001477 | 5.677977 | 33.428845 | 151.229741 | 0.032024 | 0.001026 | 0.243851 |
| 4 | Image5 | 0 | 7.315231 | 729.540579 | 27.010009 | 0.146761 | 4.283221 | 19.079108 | 174.988756 | 0.343849 | 0.118232 | 0.501140 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 3757 | Image3758 | 0 | 21.234512 | 1208.850174 | 34.768523 | 0.063774 | 2.082079 | 4.647310 | 158.437600 | 0.220666 | 0.048693 | 0.487131 |
| 3758 | Image3759 | 0 | 20.435349 | 1227.151440 | 35.030721 | 0.066763 | 2.144625 | 4.882034 | 161.158675 | 0.225931 | 0.051045 | 0.502712 |
| 3759 | Image3760 | 0 | 18.011520 | 1151.582765 | 33.934978 | 0.068396 | 2.308349 | 5.579498 | 167.130118 | 0.228930 | 0.052409 | 0.492269 |
| 3760 | Image3761 | 0 | 13.330429 | 945.732779 | 30.752769 | 0.087872 | 2.732822 | 7.757570 | 223.812932 | 0.261527 | 0.068397 | 0.480064 |
| 3761 | Image3762 | 0 | 6.110138 | 480.884025 | 21.929068 | 0.118171 | 4.110669 | 17.538826 | 239.251388 | 0.306224 | 0.093773 | 0.494333 |


3762 rows × 15 columns



In [3]: `# Head (first six rows) of brain dataframe
brain_dataframe.head(6)`

Out[3]:


| | Image | Class | Mean | Variance | Standard Deviation | Entropy | Skewness | Kurtosis | Contrast | Energy | ASM | Homogeneity | Dissimi |
|---|--------|-------|----------|-------------|--------------------|----------|----------|-----------|------------|----------|----------|-------------|---------|
| 0 | Image1 | 0 | 6.535339 | 619.587845 | 24.891522 | 0.109059 | 4.276477 | 18.900575 | 98.613971 | 0.293314 | 0.086033 | 0.530941 | 4.47 |
| 1 | Image2 | 0 | 8.749969 | 805.957634 | 28.389393 | 0.266538 | 3.718116 | 14.464618 | 63.858816 | 0.475051 | 0.225674 | 0.651352 | 3.22 |
| 2 | Image3 | 1 | 7.341095 | 1143.808219 | 33.820234 | 0.001467 | 5.061750 | 26.479563 | 81.867206 | 0.031917 | 0.001019 | 0.268275 | 5.98 |
| 3 | Image4 | 1 | 5.958145 | 959.711985 | 30.979219 | 0.001477 | 5.677977 | 33.428845 | 151.229741 | 0.032024 | 0.001026 | 0.243851 | 7.70 |
| 4 | Image5 | 0 | 7.315231 | 729.540579 | 27.010009 | 0.146761 | 4.283221 | 19.079108 | 174.988756 | 0.343849 | 0.118232 | 0.501140 | 6.83 |
| 5 | Image6 | 0 | 7.524109 | 607.395258 | 24.645390 | 0.214086 | 3.729886 | 14.471736 | 105.077882 | 0.421587 | 0.177736 | 0.598169 | 4.19 |



In [4]: *# Tail (last six rows) of brain dataframe*
brain_dataframe.tail(6)

Out[4]:

| | Image | Class | Mean | Variance | Standard Deviation | Entropy | Skewness | Kurtosis | Contrast | Energy | ASM | Homogeneity |
|-------------|-----------|-------|-----------|-------------|-----------------------|----------|----------|-----------|------------|----------|----------|-------------|
| 3756 | Image3757 | 0 | 20.976822 | 1144.456066 | 33.829810 | 0.062252 | 2.106235 | 4.798339 | 166.395916 | 0.217934 | 0.047495 | 0.488449 |
| 3757 | Image3758 | 0 | 21.234512 | 1208.850174 | 34.768523 | 0.063774 | 2.082079 | 4.647310 | 158.437600 | 0.220666 | 0.048693 | 0.487131 |
| 3758 | Image3759 | 0 | 20.435349 | 1227.151440 | 35.030721 | 0.066763 | 2.144625 | 4.882034 | 161.158675 | 0.225931 | 0.051045 | 0.502712 |
| 3759 | Image3760 | 0 | 18.011520 | 1151.582765 | 33.934978 | 0.068396 | 2.308349 | 5.579498 | 167.130118 | 0.228930 | 0.052409 | 0.492269 |
| 3760 | Image3761 | 0 | 13.330429 | 945.732779 | 30.752769 | 0.087872 | 2.732822 | 7.757570 | 223.812932 | 0.261527 | 0.068397 | 0.480064 |
| 3761 | Image3762 | 0 | 6.110138 | 480.884025 | 21.929068 | 0.118171 | 4.110669 | 17.538826 | 239.251388 | 0.306224 | 0.093773 | 0.494333 |



In [5]: *# Information of brain dataframe*
brain_dataframe.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3762 entries, 0 to 3761
Data columns (total 15 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Image                 3762 non-null   object
 1   Class                 3762 non-null   int64
 2   Mean                  3762 non-null   float64
 3   Variance              3762 non-null   float64
 4   Standard Deviation    3762 non-null   float64
 5   Entropy               3762 non-null   float64
 6   Skewness              3762 non-null   float64
 7   Kurtosis              3762 non-null   float64
 8   Contrast              3762 non-null   float64
 9   Energy                3762 non-null   float64
10  ASM                   3762 non-null   float64
11  Homogeneity           3762 non-null   float64
12  Dissimilarity         3762 non-null   float64
13  Correlation           3762 non-null   float64
14  Coarseness            3762 non-null   float64
dtypes: float64(13), int64(1), object(1)
memory usage: 441.0+ KB

```

```

In [6]: # show image columns
        print(brain_dataframe['Image'])

```

```

0      Image1
1      Image2
2      Image3
3      Image4
4      Image5
...
3757  Image3758
3758  Image3759
3759  Image3760
3760  Image3761
3761  Image3762
Name: Image, Length: 3762, dtype: object

```

image column is object dtype. contains only images name. if we drop that column it will not impact dataframe or result

```
In [7]: # drop Image column.
brain_df2 = brain_dataframe.drop(['Image'], axis = 1)
brain_df2
```

Out[7]:

| | Class | Mean | Variance | Standard Deviation | Entropy | Skewness | Kurtosis | Contrast | Energy | ASM | Homogeneity | Dissimilarity |
|-------------|-------|-----------|-------------|-----------------------|----------|----------|-----------|------------|----------|----------|-------------|---------------|
| 0 | 0 | 6.535339 | 619.587845 | 24.891522 | 0.109059 | 4.276477 | 18.900575 | 98.613971 | 0.293314 | 0.086033 | 0.530941 | 4.473340 |
| 1 | 0 | 8.749969 | 805.957634 | 28.389393 | 0.266538 | 3.718116 | 14.464618 | 63.858816 | 0.475051 | 0.225674 | 0.651352 | 3.220070 |
| 2 | 1 | 7.341095 | 1143.808219 | 33.820234 | 0.001467 | 5.061750 | 26.479563 | 81.867206 | 0.031917 | 0.001019 | 0.268275 | 5.981800 |
| 3 | 1 | 5.958145 | 959.711985 | 30.979219 | 0.001477 | 5.677977 | 33.428845 | 151.229741 | 0.032024 | 0.001026 | 0.243851 | 7.700910 |
| 4 | 0 | 7.315231 | 729.540579 | 27.010009 | 0.146761 | 4.283221 | 19.079108 | 174.988756 | 0.343849 | 0.118232 | 0.501140 | 6.834680 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 3757 | 0 | 21.234512 | 1208.850174 | 34.768523 | 0.063774 | 2.082079 | 4.647310 | 158.437600 | 0.220666 | 0.048693 | 0.487131 | 5.211730 |
| 3758 | 0 | 20.435349 | 1227.151440 | 35.030721 | 0.066763 | 2.144625 | 4.882034 | 161.158675 | 0.225931 | 0.051045 | 0.502712 | 5.083120 |
| 3759 | 0 | 18.011520 | 1151.582765 | 33.934978 | 0.068396 | 2.308349 | 5.579498 | 167.130118 | 0.228930 | 0.052409 | 0.492269 | 5.103700 |
| 3760 | 0 | 13.330429 | 945.732779 | 30.752769 | 0.087872 | 2.732822 | 7.757570 | 223.812932 | 0.261527 | 0.068397 | 0.480064 | 6.439780 |
| 3761 | 0 | 6.110138 | 480.884025 | 21.929068 | 0.118171 | 4.110669 | 17.538826 | 239.251388 | 0.306224 | 0.093773 | 0.494333 | 6.787320 |


3762 rows × 14 columns



```
In [8]: # Numerical distribution of data  
brain_df2.describe()
```

Out[8]:

| | Class | Mean | Variance | Standard Deviation | Entropy | Skewness | Kurtosis | Contrast | Energy | ASM |
|--------------|-------------|-------------|-------------|-----------------------|-------------|-------------|-------------|-------------|-------------|-------------|
| count | 3762.000000 | 3762.000000 | 3762.000000 | 3762.000000 | 3762.000000 | 3762.000000 | 3762.000000 | 3762.000000 | 3762.000000 | 3762.000000 |
| mean | 0.447368 | 9.488890 | 711.101063 | 25.182271 | 0.073603 | 4.102727 | 24.389071 | 127.961459 | 0.204705 | 0.058632 |
| std | 0.497288 | 5.728022 | 467.466896 | 8.773526 | 0.070269 | 2.560940 | 56.434747 | 109.499601 | 0.129352 | 0.058300 |
| min | 0.000000 | 0.078659 | 3.145628 | 1.773592 | 0.000882 | 1.886014 | 3.942402 | 3.194733 | 0.024731 | 0.000612 |
| 25% | 0.000000 | 4.982395 | 363.225459 | 19.058475 | 0.006856 | 2.620203 | 7.252852 | 72.125208 | 0.069617 | 0.004847 |
| 50% | 0.000000 | 8.477531 | 622.580417 | 24.951560 | 0.066628 | 3.422210 | 12.359088 | 106.737418 | 0.225496 | 0.050849 |
| 75% | 1.000000 | 13.212723 | 966.954319 | 31.095889 | 0.113284 | 4.651737 | 22.640304 | 161.059006 | 0.298901 | 0.089342 |
| max | 1.000000 | 33.239975 | 2910.581879 | 53.949809 | 0.394539 | 36.931294 | 1371.640060 | 3382.574163 | 0.589682 | 0.347725 |



```
In [9]: # check sum of null values in each columns  
brain_df2.isnull().sum()
```

```
Out[9]: Class          0
        Mean          0
        Variance      0
        Standard Deviation 0
        Entropy       0
        Skewness      0
        Kurtosis      0
        Contrast      0
        Energy        0
        ASM           0
        Homogeneity   0
        Dissimilarity 0
        Correlation   0
        Coarseness    0
        dtype: int64
```

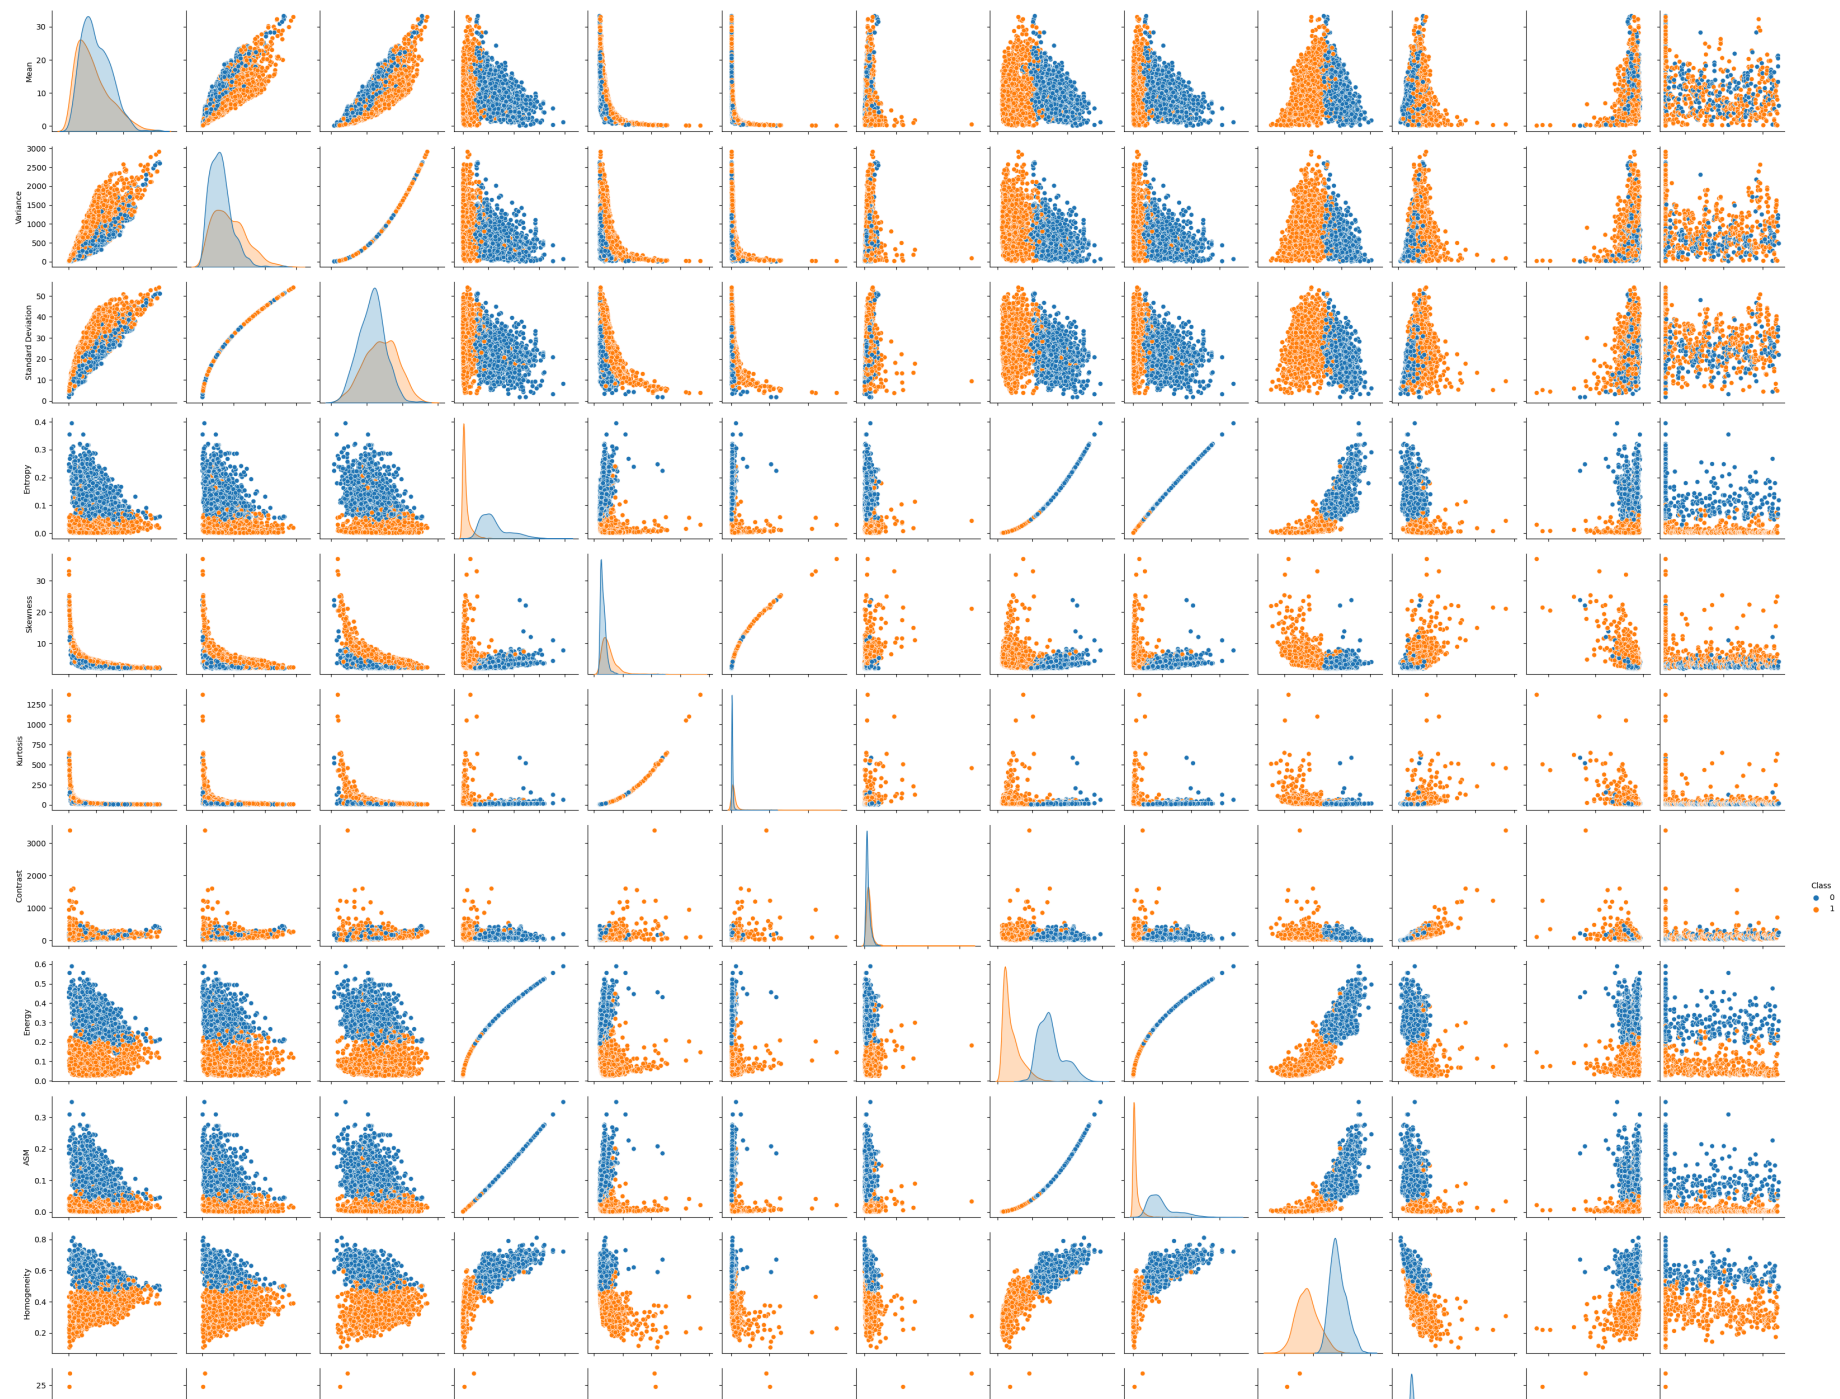
Data Visualization


```
In [10]: # Pairplot of brain dataframe
sns.pairplot(brain_df2, hue = 'Class')
```



```
C:\Users\Umaisr Ali\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
    with pd.option_context('mode.use_inf_as_na', True):  
C:\Users\Umaisr Ali\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
    with pd.option_context('mode.use_inf_as_na', True):  
C:\Users\Umaisr Ali\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
    with pd.option_context('mode.use_inf_as_na', True):  
C:\Users\Umaisr Ali\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
    with pd.option_context('mode.use_inf_as_na', True):  
C:\Users\Umaisr Ali\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
    with pd.option_context('mode.use_inf_as_na', True):  
C:\Users\Umaisr Ali\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
    with pd.option_context('mode.use_inf_as_na', True):  
C:\Users\Umaisr Ali\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
    with pd.option_context('mode.use_inf_as_na', True):  
C:\Users\Umaisr Ali\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
    with pd.option_context('mode.use_inf_as_na', True):  
C:\Users\Umaisr Ali\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
    with pd.option_context('mode.use_inf_as_na', True):  
C:\Users\Umaisr Ali\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
    with pd.option_context('mode.use_inf_as_na', True):  
C:\Users\Umaisr Ali\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
    with pd.option_context('mode.use_inf_as_na', True):
```


```
Out[10]: <seaborn.axisgrid.PairGrid at 0x21f83af1dd0>
```





```
In [11]: # Count the class columns
# 0 = Non tumor / no cancer
# 1 = Tumor / has cancer
sns.countplot(x=brain_df2["Class"])
```


```
Out[11]: <Axes: xlabel='Class', ylabel='count'>
```



Heatmap

```
In [12]: # heatmap of DataFrame  
plt.figure(figsize=(16,9))  
sns.heatmap(brain_df2)
```

```
Out[12]: <Axes: >
```



Heatmap of a correlation Matrix

```
In [13]: # correlation matrix  
brain_df2.corr()
```

Out[13]:

| | Class | Mean | Variance | Standard Deviation | Entropy | Skewness | Kurtosis | Contrast | Energy | ASM | Homogeneity | |
|-----------------------|-----------|-----------|-----------|-----------------------|-----------|-----------|-----------|-----------|-----------|-----------|-------------|-----|
| Class | 1.000000 | -0.095729 | 0.308818 | 0.285568 | -0.778180 | 0.402644 | 0.239844 | 0.212643 | -0.862413 | -0.758255 | -0.847529 | |
| Mean | -0.095729 | 1.000000 | 0.783027 | 0.790984 | -0.099729 | -0.601593 | -0.358163 | -0.050974 | -0.014863 | -0.109393 | 0.095556 | |
| Variance | 0.308818 | 0.783027 | 1.000000 | 0.975699 | -0.344432 | -0.347399 | -0.248312 | 0.135494 | -0.335470 | -0.341061 | -0.290527 | |
| Standard Deviation | 0.285568 | 0.790984 | 0.975699 | 1.000000 | -0.345127 | -0.425428 | -0.329798 | 0.117981 | -0.331103 | -0.342530 | -0.288801 | |
| Entropy | -0.778180 | -0.099729 | -0.344432 | -0.345127 | 1.000000 | -0.222222 | -0.140125 | -0.140769 | 0.971260 | 0.999213 | 0.852019 | |
| Skewness | 0.402644 | -0.601593 | -0.347399 | -0.425428 | -0.222222 | 1.000000 | 0.899713 | 0.349856 | -0.295413 | -0.209289 | -0.470054 | |
| Kurtosis | 0.239844 | -0.358163 | -0.248312 | -0.329798 | -0.140125 | 0.899713 | 1.000000 | 0.296664 | -0.172454 | -0.133741 | -0.307314 | |
| Contrast | 0.212643 | -0.050974 | 0.135494 | 0.117981 | -0.140769 | 0.349856 | 0.296664 | 1.000000 | -0.130708 | -0.139276 | -0.270119 | |
| Energy | -0.862413 | -0.014863 | -0.335470 | -0.331103 | 0.971260 | -0.295413 | -0.172454 | -0.130708 | 1.000000 | 0.961628 | 0.915988 | |
| ASM | -0.758255 | -0.109393 | -0.341061 | -0.342530 | 0.999213 | -0.209289 | -0.133741 | -0.139276 | 0.961628 | 1.000000 | 0.837139 | |
| Homogeneity | -0.847529 | 0.095556 | -0.290527 | -0.288801 | 0.852019 | -0.470054 | -0.307314 | -0.270119 | 0.915988 | 0.837139 | 1.000000 | |
| Dissimilarity | 0.556319 | -0.113864 | 0.235487 | 0.224773 | -0.502363 | 0.511931 | 0.375939 | 0.761497 | -0.545774 | -0.491813 | -0.746675 | |
| Correlation | -0.108601 | 0.293693 | 0.288037 | 0.354161 | 0.122080 | -0.570919 | -0.589211 | -0.427443 | 0.123680 | 0.121054 | 0.198639 | |
| Coarseness | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN |




we checked in information 'Coarseness' column is float datatype. contains 0 values which is not empty..... but in corelation matrix it contain NaN values..... we have to remove that column. it will not impact the result..... lets confirm in Heatmap and Barplot of Correlation Matrix

```
In [14]: # Heatmap of Correlation matrix of Brain DataFrame
plt.figure(figsize=(10,8))
sns.heatmap(brain_df2.corr(), annot = True, cmap = 'coolwarm', linewidths=2)
```

```
C:\Users\Umais Ali\anaconda3\Lib\site-packages\seaborn\matrix.py:260: FutureWarning: Format strings passed to MaskedConstant are ignored, but in future may error or produce different behavior
  annotation = ("{" + self.fmt + "}").format(val)
```

```
Out[14]: <Axes: >
```



Standard D
Sk
K
Homo
Dissi
Cor
Coa

```
In [15]: # drop coarseness column from brain_df2
brain_df3 = brain_df2.drop(['Coarseness'], axis=1)
```

Split Dataframe in Train and Test

```
In [16]: # drop dependent (Class) column, it will assign to y
X = brain_df3.drop(['Class'], axis = 1)
X.head(6)
```

Out[16]:

| | Mean | Variance | Standard Deviation | Entropy | Skewness | Kurtosis | Contrast | Energy | ASM | Homogeneity | Dissimilarity | Correlation |
|---|----------|-------------|--------------------|----------|----------|-----------|------------|----------|----------|-------------|---------------|-------------|
| 0 | 6.535339 | 619.587845 | 24.891522 | 0.109059 | 4.276477 | 18.900575 | 98.613971 | 0.293314 | 0.086033 | 0.530941 | 4.473346 | 0.9819 |
| 1 | 8.749969 | 805.957634 | 28.389393 | 0.266538 | 3.718116 | 14.464618 | 63.858816 | 0.475051 | 0.225674 | 0.651352 | 3.220072 | 0.9888 |
| 2 | 7.341095 | 1143.808219 | 33.820234 | 0.001467 | 5.061750 | 26.479563 | 81.867206 | 0.031917 | 0.001019 | 0.268275 | 5.981800 | 0.9780 |
| 3 | 5.958145 | 959.711985 | 30.979219 | 0.001477 | 5.677977 | 33.428845 | 151.229741 | 0.032024 | 0.001026 | 0.243851 | 7.700919 | 0.9641 |
| 4 | 7.315231 | 729.540579 | 27.010009 | 0.146761 | 4.283221 | 19.079108 | 174.988756 | 0.343849 | 0.118232 | 0.501140 | 6.834689 | 0.9727 |
| 5 | 7.524109 | 607.395258 | 24.645390 | 0.214086 | 3.729886 | 14.471736 | 105.077882 | 0.421587 | 0.177736 | 0.598169 | 4.193146 | 0.9764 |

```
In [17]: # output variable
# assign Class column to y
y = brain_df3['Class']
y.head(6)
```

```
Out[17]: 0    0
         1    0
         2    1
         3    1
         4    0
         5    0
         Name: Class, dtype: int64
```

```
In [18]: # split dataset into train and test
         from sklearn.model_selection import train_test_split
         X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state= 5)
```

```
In [19]: X_train
```

```
Out[19]:
```

| | Mean | Variance | Standard Deviation | Entropy | Skewness | Kurtosis | Contrast | Energy | ASM | Homogeneity | Dissimilarity | Corr |
|-------------|-----------|-------------|-----------------------|----------|----------|-----------|------------|----------|----------|-------------|---------------|------|
| 2021 | 7.885529 | 681.803866 | 26.111374 | 0.133134 | 3.806777 | 14.939540 | 277.900762 | 0.326207 | 0.106411 | 0.546301 | 6.445886 | 0 |
| 1286 | 6.639282 | 207.782336 | 14.414657 | 0.102582 | 2.607407 | 7.151179 | 33.038741 | 0.282621 | 0.079875 | 0.638838 | 2.064030 | 0 |
| 1106 | 3.186020 | 314.159477 | 17.724544 | 0.179837 | 6.354288 | 43.889536 | 165.700313 | 0.383583 | 0.147136 | 0.569095 | 4.329943 | 0 |
| 3688 | 7.752167 | 850.780349 | 29.168139 | 0.000954 | 4.667157 | 23.797444 | 162.713111 | 0.025726 | 0.000662 | 0.255565 | 7.630112 | 0 |
| 1781 | 15.907272 | 859.296845 | 29.313765 | 0.093459 | 2.257804 | 5.399814 | 111.219504 | 0.269897 | 0.072845 | 0.543017 | 3.945097 | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 3190 | 19.755035 | 1511.266702 | 38.875014 | 0.098117 | 2.359856 | 5.823204 | 155.330111 | 0.276962 | 0.076708 | 0.564342 | 4.108485 | 0 |
| 3046 | 17.384506 | 1151.899226 | 33.939641 | 0.007915 | 2.516995 | 6.645917 | 119.718675 | 0.075044 | 0.005632 | 0.381815 | 5.310392 | 0 |
| 1725 | 9.873734 | 647.095405 | 25.438070 | 0.203971 | 2.936003 | 8.936402 | 116.379595 | 0.410201 | 0.168265 | 0.653499 | 3.168139 | 0 |
| 2254 | 4.993439 | 837.608629 | 28.941469 | 0.001184 | 6.201616 | 39.682251 | 98.674332 | 0.028656 | 0.000821 | 0.218662 | 6.251872 | 0 |
| 2915 | 14.946762 | 1063.536983 | 32.611915 | 0.144926 | 2.589608 | 7.070669 | 140.031810 | 0.341421 | 0.116569 | 0.579450 | 4.088308 | 0 |

3009 rows × 12 columns



```
In [20]: X_test
```

Out[20]:

| | Mean | Variance | Standard Deviation | Entropy | Skewness | Kurtosis | Contrast | Energy | ASM | Homogeneity | Dissimilarity | Corre |
|------|-----------|------------|--------------------|----------|----------|-----------|------------|----------|----------|-------------|---------------|-------|
| 1829 | 14.244751 | 677.906461 | 26.036637 | 0.064730 | 2.245463 | 5.280704 | 68.194040 | 0.221802 | 0.049196 | 0.573020 | 3.171980 | 0.9 |
| 142 | 6.439590 | 711.402252 | 26.672125 | 0.003683 | 4.945783 | 26.258025 | 191.833333 | 0.050998 | 0.002601 | 0.252045 | 8.161609 | 0.9 |
| 2934 | 18.273666 | 960.554970 | 30.992821 | 0.053945 | 2.113554 | 4.673078 | 135.033939 | 0.201307 | 0.040524 | 0.550705 | 3.961433 | 0.9 |
| 1648 | 8.556091 | 557.147920 | 23.603981 | 0.187408 | 3.215624 | 10.791892 | 87.087468 | 0.392090 | 0.153735 | 0.626582 | 3.201900 | 0.9 |
| 1178 | 6.661285 | 761.907895 | 27.602679 | 0.004002 | 4.637791 | 22.559937 | 81.489763 | 0.052887 | 0.002797 | 0.370518 | 4.495690 | 0.9 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 3044 | 2.566971 | 254.252713 | 15.945304 | 0.001956 | 6.762827 | 47.224837 | 43.256424 | 0.036846 | 0.001358 | 0.317558 | 4.147622 | 0.9 |
| 2743 | 11.161804 | 736.299788 | 27.134845 | 0.071535 | 2.928248 | 8.891376 | 147.695042 | 0.234516 | 0.054998 | 0.510246 | 5.120198 | 0.9 |
| 3413 | 7.838409 | 300.574737 | 17.337091 | 0.134217 | 2.797981 | 8.655661 | 51.790533 | 0.327457 | 0.107228 | 0.616866 | 2.620355 | 0.9 |
| 1619 | 4.232285 | 153.173766 | 12.376339 | 0.095976 | 3.288243 | 11.029013 | 35.134788 | 0.271702 | 0.073822 | 0.639419 | 2.224797 | 0.9 |
| 3220 | 8.853958 | 573.067789 | 23.938834 | 0.006357 | 3.228923 | 10.874134 | 90.170569 | 0.067020 | 0.004492 | 0.391093 | 4.633593 | 0.9 |

753 rows × 12 columns



```
In [21]: y_train
```

```
Out[21]: 2021    0
         1286    0
         1106    1
         3688    1
         1781    0
         ..
         3190    0
         3046    1
         1725    0
         2254    1
         2915    0
         Name: Class, Length: 3009, dtype: int64
```

```
In [22]: y_test
```

```
Out[22]: 1829    0
         142     1
         2934    0
         1648    0
         1178    1
         ..
         3044    1
         2743    0
         3413    0
         1619    0
         3220    1
         Name: Class, Length: 753, dtype: int64
```

Feature Scaling

```
In [23]: from sklearn.preprocessing import StandardScaler
         sc = StandardScaler()
         X_train_sc = sc.fit_transform(X_train)
         X_test_sc = sc.transform(X_test)
```

ML Model Building

```
In [24]: from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
```

Support Vector Classifier

Train with original data

```
In [25]: from sklearn.svm import SVC
svc_classifier = SVC()
svc_classifier.fit(X_train, y_train)
y_pred_svc = svc_classifier.predict(X_test)
accuracy_score(y_test, y_pred_svc)
```

```
Out[25]: 0.7861885790172642
```

Train with scaled data

```
In [26]: svc_classifier2 = SVC()
svc_classifier2.fit(X_train_sc, y_train)
y_pred_svc_sc = svc_classifier2.predict(X_test_sc)
accuracy_score(y_test, y_pred_svc_sc)
```

```
Out[26]: 0.9774236387782205
```

Logistic Regression

Train with original data

```
In [27]: from sklearn.linear_model import LogisticRegression
lr_classifier = LogisticRegression(random_state=51, penalty = 'l2')
lr_classifier.fit(X_train, y_train)
y_pred_lr = lr_classifier.predict(X_test)
accuracy_score(y_test, y_pred_lr)
```

```
C:\Users\Umaisr Ali\anaconda3\Lib\site-packages\sklearn\linear_model\_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

Out[27]: 0.9096945551128818

Train with scaled data

```
In [28]: lr_classifier2 = LogisticRegression(random_state=51, penalty = 'l2')  
lr_classifier2.fit(X_train_sc, y_train)  
y_pred_lr_sc = lr_classifier2.predict(X_test_sc)  
accuracy_score(y_test, y_pred_lr_sc)
```

Out[28]: 0.9787516600265604

KNN - K-Nearesr Neighbor Classifier

Train with original data

```
In [29]: from sklearn.neighbors import KNeighborsClassifier  
knn_classifier = KNeighborsClassifier(n_neighbors = 5, metric = 'minkowski', p = 2)  
knn_classifier.fit(X_train, y_train)  
y_pred_knn = knn_classifier.predict(X_test)  
accuracy_score(y_test, y_pred_knn)
```

Out[29]: 0.8061088977423638

Train with scaled data

```
In [30]: knn_classifier2 = KNeighborsClassifier(n_neighbors = 5, metric = 'minkowski', p = 2)  
knn_classifier2.fit(X_train_sc, y_train)
```

```
y_pred_knn_sc = knn_classifier2.predict(X_test_sc)
accuracy_score(y_test, y_pred_knn_sc)
```

Out[30]: 0.9774236387782205

Naive Bayes Classifier

Train with original data

```
In [31]: from sklearn.naive_bayes import GaussianNB
nb_classifier = GaussianNB()
nb_classifier.fit(X_train, y_train)
y_pred_nb = nb_classifier.predict(X_test)
accuracy_score(y_test, y_pred_nb)
```

Out[31]: 0.953519256308101

Train with scaled data

```
In [32]: nb_classifier2 = GaussianNB()
nb_classifier2.fit(X_train_sc, y_train)
y_pred_nb_sc = nb_classifier2.predict(X_test_sc)
accuracy_score(y_test, y_pred_nb_sc)
```

Out[32]: 0.9548472775564409

Decision Tree Classifier

Train with original data

```
In [33]: from sklearn.tree import DecisionTreeClassifier
dt_classifier = DecisionTreeClassifier(criterion = 'entropy', random_state = 51)
dt_classifier.fit(X_train, y_train)
y_pred_dt = dt_classifier.predict(X_test)
accuracy_score(y_test, y_pred_dt)
```

Out[33]: 0.9814077025232404

Train with scaled data

```
In [34]: dt_classifier2 = DecisionTreeClassifier(criterion = 'entropy', random_state = 51)
dt_classifier2.fit(X_train_sc, y_train)
y_pred_dt_sc = dt_classifier2.predict(X_test_sc)
accuracy_score(y_test, y_pred_dt_sc)
```

Out[34]: 0.9814077025232404

Random Forest Classifier

Train with original data

```
In [35]: from sklearn.ensemble import RandomForestClassifier
rf_classifier = RandomForestClassifier(n_estimators = 20, criterion = 'entropy', random_state = 51)
rf_classifier.fit(X_train, y_train)
y_pred_rf = rf_classifier.predict(X_test)
accuracy_score(y_test, y_pred_rf)
```

Out[35]: 0.9827357237715804


Train with scaled data

```
In [36]: rf_classifier2 = RandomForestClassifier(n_estimators = 20, criterion = 'entropy', random_state = 51)
rf_classifier2.fit(X_train_sc, y_train)
y_pred_rf_sc = rf_classifier2.predict(X_test_sc)
accuracy_score(y_test, y_pred_rf_sc)
```

Out[36]: 0.9827357237715804

Confusion Matrix


```
In [37]: cm = confusion_matrix(y_test, y_pred_rf_sc)
plt.title('Heatmap of Confusion Matrix', fontsize = 15)
sns.heatmap(cm, annot = True)
plt.show()
```



Classification Report Of model

```
In [38]: print(classification_report(y_test, y_pred_rf_sc))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.98 | 0.99 | 0.98 | 411 |
| 1 | 0.99 | 0.97 | 0.98 | 342 |
| accuracy | | | 0.98 | 753 |
| macro avg | 0.98 | 0.98 | 0.98 | 753 |
| weighted avg | 0.98 | 0.98 | 0.98 | 753 |

Cross-validation of the ML model

```
In [39]: # Cross validation
from sklearn.model_selection import cross_val_score
cross_validation = cross_val_score(estimator = rf_classifier2, X = X_train_sc, y = y_train, cv = 10)
print("Cross validation accuracy of SVC model = ", cross_validation)
print("\nCross validation mean accuracy of SVC model = ", cross_validation.mean())
```

```
Cross validation accuracy of SVC model = [0.98671096 0.99335548 0.99003322 0.97009967 0.9833887 0.97674419
0.99335548 0.99003322 0.99335548 0.99666667]
```

```
Cross validation mean accuracy of SVC model = 0.9873743078626799
```

Test Model

```
In [45]: # we have one patient data
patient1 = [17.99,10.38,122.8,1001.0,0.1184,0.2776,0.3001,
0.1471,0.2419,0.07871,1.095,0.9053]
```

```
In [41]: # convert list into array and scale
patient1_sc = sc.transform(np.array([patient1]))
patient1_sc
```

```
C:\Users\Umar Ali\anaconda3\Lib\site-packages\sklearn\base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
warnings.warn(
```

```
Out[41]: array([[ 1.49256979e+00, -1.49381108e+00,  1.11029758e+01,
                  1.40960224e+04, -1.54791906e+00, -4.16334494e-01,
                  -1.13805522e+00, -4.48064734e-01,  3.09915646e+00,
                  -3.10934480e+00, -1.93429813e+00, -1.84574993e+00]])
```

```
In [42]: # predict patient1_sc scale data
# zero mean malignant patient has cancer
predict = rf_classifier2.predict(patient1_sc)
predict
```

```
Out[42]: array([0], dtype=int64)
```

```
In [43]: # write if else statement to print result in clear format
if predict[0] == 0:
    print ('Patient has *** NO *** Tumor / Cancer')
else:
    print ('Patient *** HAS *** Tumor / Cancer')
```

Patient has *** NO *** Tumor / Cancer

```
In [44]: # confusion matrix
print('Confusion matrix of Random Forest model: \n',confusion_matrix(y_test, y_pred_rf_sc),'\n')

# show the accuracy
print('Accuracy of Random Forest model = ',accuracy_score(y_test, y_pred_rf_sc))
```

Confusion matrix of Random Forest model:

```
[[408   3]
 [ 10 332]]
```

Accuracy of Random Forest model = 0.9827357237715804